

# In-Class Problem Set: Exploring Movie Data with Distribution and Color (R + GitHub *or* Canvas)

**Goal.** Use real movie data to practice visualizing distributions, comparing groups, and encoding multiple variables in a single plot. You will obtain the dataset from the course materials (GitHub *or* Canvas), build a reproducible workflow, generate several plots, interpret what they show, and submit your work via **GitHub or Canvas**.

## What to submit (GitHub *or* Canvas).

- A script file: `scripts/lab.R`
- A short write-up: `outputs/writeup.md`
- Saved figures in `figures/` (see requirements below)

If you submit via Canvas, upload the same files listed above as individual files (or upload a single zip that preserves the directory structure).

## Rules.

- Work inside an **R Project**.
- Use a **sequential, hard-coded workflow** (no user-defined functions).
- Save figures using code (`ggsave`); do not use screenshots.
- If you choose the **GitHub submission option**, Git commands must be run in the **Terminal tab**, not the R Console.

## Submission options

You may submit this assignment using **either** of the following methods:

- **GitHub submission (recommended):** Commit and push your work to your GitHub repository. You will include Git proof (`git status` and `git log`) in your write-up.
- **Canvas submission:** Upload the required files directly to Canvas. You do *not* need to use GitHub if you choose this option.

Both submission methods are graded using the same rubric.

## Mini codebook (use this; do not guess)

Each row represents one movie. Relevant variables include:

- **budget:** Production budget in USD (0 if unavailable).
- **revenue:** Worldwide box office revenue in USD (0 if unavailable).
- **director:** Director of the movie.
- **runtime:** Movie length in minutes.
- **vote\_average:** Average user rating (0–10).
- **vote\_count:** Number of user votes.
- **popularity:** Popularity score based on user engagement.

## Questions

### 1. Get the movie dataset (proof required).

(a) Choose **one** method:

- **GitHub option:** In the **Terminal tab**, run:

```
git status  
git pull
```

- **Canvas option:** Download the movie dataset from Canvas and place it in your project **data/** folder.

(b) **Pseudo-code (follow structure; fill in details):**

```
# confirm working directory  
-----()  
  
# list files in data directory  
-----("data")  
  
# visually confirm expected movie file exists
```

(c) Confirm the movie dataset exists in your project.

(d) **Proof (write-up):** In **outputs/writeup.md**, paste:

- the output of `getwd()`,
- the output of `list.files("data")`.

### 2. Load and summarize the dataset.

(a) Load the movie dataset into an object named `df`.

(b) **Pseudo-code (intentionally incomplete):**

```
df <- read.csv("data/_____.csv")  
  
# quick structure checks (choose at least two)  
-----(df)  
-----(df)  
-----(df)
```

(c) Summarize the dataset to understand its structure.

(d) **Proof (write-up):** Report:

- number of rows and columns,
- the range of `budget`,
- the range of `revenue`.

### 3. Plot distributions: movie budget and revenue.

Create two histograms:

- one for `budget`,
- one for `revenue`.
- Repeat the process again and chose alternative bin sizes to assess how that changes the interpretation. Use one of the following formulas:

– **Freedman–Diaconis (bin width):**

$$h = 2 \cdot \text{IQR}(x) \cdot n^{-1/3}$$

– **Scott's Rule (bin width):**

$$h = 3.5 \cdot s(x) \cdot n^{-1/3}$$

- Sturges' Rule (number of bins):

$$k = \lceil \log_2(n) + 1 \rceil$$

#### 4. Pseudo-code (structure only):

```
# budget histogram
ggplot(df, aes(x = _____)) +
  geom_histogram(_____)

# revenue histogram
ggplot(df, aes(x = _____)) +
  geom_histogram(_____)
```

Save the plots as:

- figures/budget\_hist.png
- figures/revenue\_hist.png

#### 5. Identify top-grossing directors and compare revenue.

- Identify the **top three directors** by total box office revenue.
- Subset the data to movies directed by these three directors.
- Create a **boxplot** showing the distribution of **revenue** for each director.

#### 6. Pseudo-code (leave blanks):

```
top_directors <- df %>%
  group_by(_____) %>%
  summarize(total_revenue = _____) %>%
  arrange(_____) %>%
  slice(___)

df_top <- df %>%
  filter(_____%in% top_directors$_____)
```

Save the plot as:

figures/revenue\_by\_director.png

#### 7. Scatter plot with size and color encodings.

Create a scatter plot with:

- x-axis: **budget**
- y-axis: **revenue**

Then:

- map one quantitative variable to **point size**,
- map one categorical variable to **color**.

#### 8. Pseudo-code (structure only):

```
ggplot(df, aes(x = _____, y = _____,
                size = _____, color = _____)) +
  geom_point(_____)
```

Save the plot as:

figures/budget\_revenue\_scatter.png

## 9. Interpretation (write-up required).

In `outputs/writeup.md`, write 10–14 sentences addressing:

- what the budget and revenue distributions reveal,
- how revenues differ across top directors,
- what relationships are most salient in the scatter plot,
- how size and color encodings affect interpretation.

## 10. Submit your work (GitHub or Canvas) + proof required.

### (a) Choose ONE submission method:

- **GitHub option:** Commit and push your work to GitHub.
- **Canvas option:** Upload `scripts/lab.R`, `outputs/writeup.md`, and all required figures to Canvas.

### (b) Pseudo-code (submission skeleton):

```
# GitHub option
git status
git add -----
git commit -m "-----"
git push
```

### (c) Proof (write-up):

- If using **GitHub**: paste `git status` and `git log -1`.
- If using **Canvas**: paste `list.files("scripts")`, `list.files("outputs")`, and `list.files("figures")`, and state that you submitted via Canvas.

## Optional challenge (if you finish early)

Choose one plot and create an alternative version optimized for a **general public** audience. In 5–7 sentences, explain what design choices you changed and why.

## Checklist (before you leave)

- `scripts/lab.R` runs top-to-bottom
- `outputs/writeup.md` exists and includes interpretation + proofs
- Required figures exist in `figures/`
- Work is either committed and pushed to GitHub or uploaded to Canvas